

Predicting heart disease using machine learning

This notebook looks into using various python-based learning and data science libraries in an attempt to build a machine learning capable of predicting whether or not someone has heart disease based on their medical attributes.

we are going to take the following approach:

1. Problem definition
2. data
3. Evaluation
4. Features
5. Modelling
6. Experimentation

1. Problem definition

In a statement,

Given clinical parameters about a patient, can we predict whether or not they have heart disease.

2. Data

The original data come from the Cleveland data from the learning repository.

<http://archive.ics.edu/ml/datasets/heart+Disease> (<http://archive.ics.edu/ml/datasets/heart+Disease>)

There is also a version of it available on Kaggle.

<https://www.kaggle.com/ronitf/heart-disease-uci> (<https://www.kaggle.com/ronitf/heart-disease-uci>)

- Always see important attributes from the given data and select only those which will be useful.
- The original heart disease dataset contains 76 attributes. but, only 14 are selected. We should understand meaning and dependency of attributes on each other.
- Get all data in useful format.

3. Evaluation

If we can reach 95% accuracy at predicting whether or not a patient has heart disease during the proof of concept, we will pursue the project.

- Always decide desired accuracy vs severity of the model.(or customer requirement)

4. Features

Create data dictionary This is where you will get different information about each of the features in your data. You can do this via doing own research (such as looking at the links above) or by talking to a subject matter expert(Someone who knows about the dataset).

Create data dictionary

1. age - age in years
2. sex - (1 = male; 0 = female)
3. cp - chest pain type
4. trestbps - resting blood pressure (in mm Hg on admission to the hospital)
4. chol - serum cholestoral in mg/dl
5. fbs - (fasting blood sugar > 120 mg/dl) (1 = true; 0 = false)
6. restecg - resting electrocardiographic results
7. thalach - maximum heart rate achieved
8. exang - exercise induced angina (1 = yes; 0 = no)
9. oldpeak - ST depression induced by exercise relative to rest
10. slope - the slope of the peak exercise ST segment
11. ca - number of major vessels (0-3) colored by flourosopy
12. thal - 3 = normal; 6 = fixed defect; 7 = reversable defect
13. target - 1 or 0

In []: 1

Preparing the tools

we are going to use pandas, matplotlib and NumPy for data analysis and manipulation.

```
In [1]: 1 # Import all tools we need
        2 # Regular EDA(Exploratory data analysis)
        3 import numpy as np
        4 import pandas as pd
        5 import matplotlib.pyplot as plt
        6 import seaborn as sns
        7 %matplotlib inline
        8 # Models from Scikit-Learn
        9 from sklearn.linear_model import LogisticRegression
       10 from sklearn.neighbors import KNeighborsClassifier
       11 from sklearn.ensemble import RandomForestClassifier
       12
       13 # Moel Evaluation
       14 from sklearn.model_selection import train_test_split, cross_val_score
       15 from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
       16 from sklearn.metrics import confusion_matrix, classification_report
       17 from sklearn.metrics import precision_score, recall_score, f1_score
       18 from sklearn.metrics import plot_roc_curve
```

LOAD DATA

```
In [2]: 1 df = pd.read_csv("heart_disease.csv")
        2 df.head()
```

```
Out[2]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

Data Exploration(Exploratory data analysis or EDA)

the goal here is to find out more about the data and become a subject matter expert on the dataset you are working with.

1. What questions are you trying to solve?
2. What kind of data so we have and how so we treat different types?
3. What is missing from the data and how do you deal with it?
4. Where are the outliers and why should you care about them?
5. How can you add, change or remove features to get more out of your data?

```
In [3]: 1 df.head()
```

```
Out[3]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

```
In [4]: 1 df.tail()
```

```
Out[4]:
```


	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	0
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	0
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0

```
In [5]: 1 df["target"].value_counts()
```

```
Out[5]: 1    165
        0    138
        Name: target, dtype: int64
```

```
In [6]: 1 df["target"].value_counts().plot(kind="bar", color=["salmon", "skyblue"])
```

```
Out[6]: <matplotlib.axes._subplots.AxesSubplot at 0x136a80a7580>
```



Visualizing the counts of 0 and 1 are nearly equal. so, problem is balanced classification problem.

In [7]:

1

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         303 non-null    int64
 1   sex         303 non-null    int64
 2   cp          303 non-null    int64
 3   trestbps    303 non-null    int64
 4   chol        303 non-null    int64
 5   fbs         303 non-null    int64
 6   restecg     303 non-null    int64
 7   thalach     303 non-null    int64
 8   exang       303 non-null    int64
 9   oldpeak     303 non-null    float64
10   slope       303 non-null    int64
11   ca          303 non-null    int64
12   thal        303 non-null    int64
13   target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

In [8]:

1

df.isna().sum()

```
Out[8]: age         0
sex         0
cp          0
trestbps    0
chol        0
fbs         0
restecg     0
thalach     0
exang       0
oldpeak     0
slope       0
ca          0
thal        0
target      0
dtype: int64
```

In [9]:

1

df.describe()

```
Out[9]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalact
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053	149.646865
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860	22.905161
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.500000
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	153.000000
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000	166.000000
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000

Heart disease frequency according to sex


```
In [10]: 1 df.sex.value_counts()
```

```
Out[10]: 1    207
0     96
Name: sex, dtype: int64
```

```
In [11]: 1 # Compare Sex Vs Target
2 pd.crosstab(df.target, df.sex)
```

```
Out[11]:    sex  0   1
target
0      24  114
1      72   93
```

```
In [12]: 1 pd.crosstab(df.target, df.sex).plot(kind="bar",
2                                             figsize=(10,6),
3                                             color=["skyblue", "red"])
4 plt.title("Heart disease frequency for sex")
5 plt.xlabel("0= No disease, 1=Heart disease")
6 plt.ylabel("Amount")
7 plt.legend(["Female", "Male"])
8 plt.xticks(rotation=0);
```




```
In [13]: 1 df["thalach"].value_counts()

Out[13]: 162    11
        160     9
        163     9
        173     8
        152     8
        ..
        129     1
        128     1
        127     1
        124     1
        71      1
        Name: thalach, Length: 91, dtype: int64
```

Age Vs.thalach for heart rate


```
In [14]: 1 # Create another figure
        2 plt.figure(figsize=(10,6))
        3
        4 #scatter with positive example
        5 plt.scatter(df.age[df.target==1],
        6             df.thalach[df.target==1])
        7 #scatter with negative example
        8 plt.scatter(df.age[df.target==0],
        9             df.thalach[df.target==0])
        10
        11 plt.title("Heart disease is function of age and thalach")
        12 plt.xlabel("Age")
        13 plt.ylabel("Thalach-Max Heart rate")
        14 plt.legend(["Disease", "No Disease"])
```

```
Out[14]: <matplotlib.legend.Legend at 0x136a83d4ca0>
```




In [15]:

```
1 # Check the distribution of the age column with a histogram
2 df.age.plot.hist();
```



In [16]:

```
1 df.thalach.plot.hist();
```



In [17]:


```
1 pd.crosstab(df.cp, df.target)
```

Out[17]:

	target	
	0	1
cp		
0	104	39
1	9	41
2	18	69
3	7	16

In [18]:

```
1 pd.crosstab(df.cp, df.target).plot(kind="bar",
2                                     figsize=(10,6),
3                                     color=["skyblue", "green"])
4 plt.title("Heart disease frequency per chest pain type")
5 plt.xlabel("Chest pain Type")
6 plt.ylabel("Amount")
7 plt.legend(["No-disease", "Disease"])
8 plt.xticks(rotation=0);
```



In [19]:

```
1 df.head()
```

Out[19]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

In [20]:

```
1 # Make a correlation matrix
2 df.corr()
```


Out[20]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exa
age	1.000000	-0.098447	-0.068653	0.279351	0.213678	0.121308	-0.116211	-0.398522	0.0968
sex	-0.098447	1.000000	-0.049353	-0.056769	-0.197912	0.045032	-0.058196	-0.044020	0.1416
cp	-0.068653	-0.049353	1.000000	0.047608	-0.076904	0.094444	0.044421	0.295762	-0.3942
trestbps	0.279351	-0.056769	0.047608	1.000000	0.123174	0.177531	-0.114103	-0.046698	0.0676
chol	0.213678	-0.197912	-0.076904	0.123174	1.000000	0.013294	-0.151040	-0.009940	0.0670
fbs	0.121308	0.045032	0.094444	0.177531	0.013294	1.000000	-0.084189	-0.008567	0.0256
restecg	-0.116211	-0.058196	0.044421	-0.114103	-0.151040	-0.084189	1.000000	0.044123	-0.0707
thalach	-0.398522	-0.044020	0.295762	-0.046698	-0.009940	-0.008567	0.044123	1.000000	-0.3788
exang	0.096801	0.141664	-0.394280	0.067616	0.067023	0.025665	-0.070733	-0.378812	1.0000
oldpeak	0.210013	0.096093	-0.149230	0.193216	0.053952	0.005747	-0.058770	-0.344187	0.2882
slope	-0.168814	-0.030711	0.119717	-0.121475	-0.004038	-0.059894	0.093045	0.386784	-0.2577
ca	0.276326	0.118261	-0.181053	0.101389	0.070511	0.137979	-0.072042	-0.213177	0.1157
thal	0.068001	0.210041	-0.161736	0.062210	0.098803	-0.032019	-0.011981	-0.096439	0.2067
target	-0.225439	-0.280937	0.433798	-0.144931	-0.085239	-0.028046	0.137230	0.421741	-0.4367



In [21]:

```
1 # Let's make correlation matrix more visible
2 corr_matrix = df.corr()
3 fig, ax = plt.subplots(figsize=(15,10))
4 ax = sns.heatmap(corr_matrix,
5                  annot=True,
6                  linewidths=0.5,
7                  fmt=".2f",
8                  cmap="YlGnBu")
```



5. Modelling

In [22]:

```
1 df.head()
```

Out[22]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

In [23]:

```
1 #Splitting the data
2 x= df.drop("target", axis=1)
3 y= df["target"]
```

```
In [24]: ▶ 1 y
          2
```

```
Out[24]: 0      1
         1      1
         2      1
         3      1
         4      1
         ..
        298     0
        299     0
        300     0
        301     0
        302     0
        Name: target, Length: 303, dtype: int64
```

```
In [25]: ▶ 1 x.shape, y.shape
```

```
Out[25]: ((303, 13), (303,))
```

```
In [26]: ▶ 1 np.random.seed(55)
          2 x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.2)
```

```
In [27]: ▶ 1 x_train.shape, x_test.shape
```

```
Out[27]: ((242, 13), (61, 13))
```

Now we've split into training and test set, time to choose right estimator. we'll train it on training set and test it on test set.

We're going to try three different ML models:

1. Logistic Regression
2. K-Nearest neighbours classifier
3. Random Forest Classifier

*Logistic regression. despites its name, is a linear model for classification rather than regression. It is also known as the literature as logit regression, maximum-entropy classification or log-linear classifier.

```
In [28]: ▶ 1 # Put models in a dictionary
2
3 models = {"Logistic Regression": LogisticRegression(),
4           "KNN": KNeighborsClassifier(),
5           "Random Forest": RandomForestClassifier()}
6
7 #Create a function to fit and score models
8 def fit_and_scores (models, x_train, x_test, y_train, y_test):
9     """
10     Fits and evaluate given machine learning models.
11     models: a dict of different Scikit-Learn machine learning models
12     x_train : training data(no labels)
13     x_test : testing data(no labels)
14     y_train : training labels
15     y_test : testing labels
16     """
17
18     np.random.seed(55)
19     model_scores = {}
20
21     for name, model in models.items():
22         model.fit(x_train, y_train)
23         model_scores[name] = model.score(x_test, y_test)
24
25     return model_scores
```

```
In [29]: ▶ 1 model_scores = fit_and_scores(models=models,
2           x_train=x_train,
3           x_test=x_test,
4           y_train=y_train,
5           y_test = y_test)
6 model_scores
```

C:\Users\suhass\AI\Project2\env\lib\site-packages\sklearn\linear_model_logistic.py:938: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)


Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
(https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[29]: {'Logistic Regression': 0.8360655737704918,
'KNN': 0.7377049180327869,
'Random Forest': 0.8032786885245902}
```

In [30]: ▶

```
1 # Model Comparisons
2 model_compare = pd.DataFrame(model_scores, index=["accuracy"])
3 model_compare.T.plot.bar();
```



now we have a baseline model. and we know a model's first predictions aren't always what we should our next steps off. What should do??

Let's look at the following:

1. Hyperparameter tuning
2. Feature importance
3. Confusion matrix
4. Cross-validation
5. Precision
6. Recall
7. F1 score
8. Classification report
9. ROC Curve
10. Area under the curve (AUC)

Hyperparameter tuning

In [31]:

```
1 # Let's tune KNN
2 train_scores = []
3 test_scores = []
4
5 # create a list of range of neighbours
6 neighbors = range(1,21)
7
8 # Setup KNN instance
9 knn = KNeighborsClassifier()
10
11 # Loops through different neighbors
12 for i in neighbors :
13     knn.set_params(n_neighbors=i)
14     knn.fit(x_train, y_train)
15     train_scores.append(knn.score(x_train, y_train))
16     test_scores.append(knn.score(x_test, y_test))
```

In [32]:

```
1 train_scores
2
```

Out[32]: [1.0,
0.7975206611570248,
0.7727272727272727,
0.7479338842975206,
0.7603305785123967,
0.7603305785123967,
0.743801652892562,
0.743801652892562,
0.7396694214876033,
0.7107438016528925,
0.7024793388429752,
0.7148760330578512,
0.6859504132231405,
0.7148760330578512,
0.6900826446280992,
0.7024793388429752,
0.6776859504132231,
0.6776859504132231,
0.6818181818181818,
0.6900826446280992]


In [33]:

```
1 test_scores
```

Out[33]: [0.6885245901639344,
0.6229508196721312,
0.6721311475409836,
0.6885245901639344,
0.7377049180327869,
0.6885245901639344,
0.7213114754098361,
0.7049180327868853,
0.6885245901639344,
0.6885245901639344,
0.6885245901639344,
0.6885245901639344,
0.6885245901639344,
0.7049180327868853,
0.6885245901639344,
0.7213114754098361,
0.6721311475409836,
0.7213114754098361,
0.7049180327868853,
0.7213114754098361,
0.7213114754098361]

```
In [34]: ▶ 1 plt.plot(neighbors, train_scores, label="Train score")
2 plt.plot(neighbors, test_scores, label="Test score")
3 plt.xticks(np.arange(1,21,1))
4 plt.xlabel("Number of neighbors")
5 plt.ylabel("Model Score")
6 plt.legend()
7 print(f"Maximum KNN score on the test data:{max(test_scores)*100:.2f}%")
```

Maximum KNN score on the test data:73.77%



Hyperparameter tuning by RandomizedSearchCV

We are going to tune:

- LogisticRegression()
- RandomForestClassifier()

... using RandomizedSearchCV

```
In [46]: ▶ 1 # Create a hyperparameter grid for logistic regression
2
3 log_reg_grid = {"C": np.logspace(-4, 4, 20),
4                 "solver" : ["liblinear"]}
5 # Create a hyperparameter grid for RandomForestClassifier
6
7 rf_grid = {"n_estimators": np.arange(10, 1000, 50),
8            "max_depth": [None, 3, 5, 10],
9            "min_samples_split": np.arange(2, 20, 2),
10           "min_samples_leaf" : np.arange(1, 20, 2)}
```

now we've got hyperparameter grid setup for each of our models, let's tune them using RandomizedSearchCV

```
In [47]: ▶ 1 #Tune LogisticRegression
2 np.random.seed(77)
3
4 # Setup random hyperparameter search for LogisticRegression
5 rs_log_reg = RandomizedSearchCV(LogisticRegression(),
6                                 param_distributions=log_reg_grid,
7                                 cv=5,
8                                 n_iter = 20,
9                                 verbose =True)
10
11 #Fit the model
12 rs_log_reg.fit(x_train, y_train)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
 [Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 0.7s finished

```
Out[47]: RandomizedSearchCV(cv=5, error_score=nan,
                             estimator=LogisticRegression(C=1.0, class_weight=None,
                                                            dual=False, fit_intercept=True,
                                                            intercept_scaling=1,
                                                            l1_ratio=None, max_iter=100,
                                                            multi_class='auto', n_jobs=None,
                                                            penalty='l2', random_state=None,
                                                            solver='lbfgs', tol=0.0001,
                                                            verbose=0, warm_start=False),
                             iid='deprecated', n_iter=20, n_jobs=None,
                             param_distributions={'C':...
4.83293024e-03, 1.27427499e-02, 3.35981829e-02, 8.85866790e-02,
2.33572147e-01, 6.15848211e-01, 1.62377674e+00, 4.28133240e+00,
1.12883789e+01, 2.97635144e+01, 7.84759970e+01, 2.06913808e+02,
5.45559478e+02, 1.43844989e+03, 3.79269019e+03, 1.00000000e+04]),
                             'solver': ['liblinear']},
                             pre_dispatch='2*n_jobs', random_state=None, refit=True,
                             return_train_score=False, scoring=None, verbose=True)
```

```
In [48]: ▶ 1 rs_log_reg.best_params_
```

```
Out[48]: {'solver': 'liblinear', 'C': 0.08858667904100823}
```

```
In [49]: ▶ 1 rs_log_reg.score(x_test, y_test)
```

```
Out[49]: 0.8032786885245902
```

Now we've tuned LogisticRegression, let's do it for the RandomForestClassifier()..

In [50]: ►

```
1 #Tune LogisticRegression
2 np.random.seed(77)
3
4 # Setup random hyperparameter search for LogisticRegression
5 rs_rf = RandomizedSearchCV(RandomForestClassifier(),
6                             param_distributions= rf_grid,
7                             cv=5,
8                             n_iter = 20,
9                             verbose =True)
10
11 #Fit the model
12 rs_rf.fit(x_train, y_train)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 100 out of 100 | elapsed: 2.5min finished

Out[50]: RandomizedSearchCV(cv=5, error_score=nan,
 estimator=RandomForestClassifier(bootstrap=True,
 ccp_alpha=0.0,
 class_weight=None,
 criterion='gini',
 max_depth=None,
 max_features='auto',
 max_leaf_nodes=None,
 max_samples=None,
 min_impurity_decrease=0.0,
 min_impurity_split=None,
 min_samples_leaf=1,
 min_samples_split=2,
 min_weight_fraction_leaf=0.0,
 n_estimators=100,
 n_jobs...
 param_distributions={'max_depth': [None, 3, 5, 10],
 'min_samples_leaf': array([1, 3, 5, 7,
9, 11, 13, 15, 17, 19]),
 'min_samples_split': array([2, 4, 6,
8, 10, 12, 14, 16, 18]),
 'n_estimators': array([10, 60, 110, 160,
210, 260, 310, 360, 410, 460, 510, 560, 610,
660, 710, 760, 810, 860, 910, 960])},
 pre_dispatch='2*n_jobs', random_state=None, refit=True,
 return_train_score=False, scoring=None, verbose=True)

In [51]: ►

```
1 # Find the best hyperparameters
2 rs_rf.best_params_
```

Out[51]: {'n_estimators': 360,
 'min_samples_split': 18,
 'min_samples_leaf': 1,
 'max_depth': None}

In [53]: ►

```
1 #Evaluate the RandomizedSearch RandomForestClassifier
2
3 rs_rf.score(x_test, y_test)
```

Out[53]: 0.8360655737704918

Hyperparameter Tuning with GridSearchCV

since LogisticRegression model provides the best scores so far, we'll try and improve them again using GridSearchCV

```
In [58]: 1 #Different Hyperparameters for LogisticRegression
2 log_reg_grid = { "C": np.logspace(-4, 4, 30),
3                 "solver": ["liblinear"]}
4
5 # setup grid hyperparameters search for LogisticRegression
6 gs_log_reg = GridSearchCV(LogisticRegression(),
7                           param_grid=log_reg_grid,
8                           cv=5,
9                           verbose=True)
10
11 # Fit the hyperparameter search model
12 gs_log_reg.fit(x_train, y_train)
```

Fitting 5 folds for each of 30 candidates, totalling 150 fits

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 150 out of 150 | elapsed: 1.6s finished

```
Out[58]: GridSearchCV(cv=5, error_score=nan,
                      estimator=LogisticRegression(C=1.0, class_weight=None, dual=False,
                                                    fit_intercept=True,
                                                    intercept_scaling=1, l1_ratio=None,
                                                    max_iter=100, multi_class='auto',
                                                    n_jobs=None, penalty='l2',
                                                    random_state=None, solver='lbfgs',
                                                    tol=0.0001, verbose=0,
                                                    warm_start=False),
                      iid='deprecated', n_jobs=None,
                      param_grid={'C': array([1.00000000e-04, 1.8...
2.04335972e-01, 3.85662042e-01, 7.27895384e-01, 1.37382380e+00,
2.59294380e+00, 4.89390092e+00, 9.23670857e+00, 1.74332882e+01,
3.29034456e+01, 6.21016942e+01, 1.17210230e+02, 2.21221629e+02,
4.17531894e+02, 7.88046282e+02, 1.48735211e+03, 2.80721620e+03,
5.29831691e+03, 1.00000000e+04]),
                                'solver': ['liblinear']}),
                      pre_dispatch='2*n_jobs', refit=True, return_train_score=False,
                      scoring=None, verbose=True)
```

```
In [59]: 1 # Check the best hyperparameter
2 gs_log_reg.best_params_
3
```

```
Out[59]: {'C': 0.1082636733874054, 'solver': 'liblinear'}
```

```
In [60]: 1 # Evaluate the score
2 gs_log_reg.score(x_test, y_test)
```

```
Out[60]: 0.8032786885245902
```

```
In [61]: 1 model_scores
```

```
Out[61]: {'Logistic Regression': 0.8360655737704918,
          'KNN': 0.7377049180327869,
          'Random Forest': 0.8032786885245902}
```

Evaluating Our tuned machine learning classifier, beyond accuracy

- ROC Curve and AUC score

- Confusion Matrix
- Classification report
- Precision
- Recall
- F1

... and it would be great if cross validation was used where possible.

to make comparisons and evaluate our trained model, first we need to make predictions.

```
In [62]: 1 # Make predictions with tuned model
        2 y_preds = gs_log_reg.predict(x_test)
```

```
In [63]: 1 y_preds
```


```
Out[63]: array([1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1,
                1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0,
                0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1], dtype=int64)
```

```
In [64]: 1 y_test
```

```
Out[64]: 92      1
        121      1
        53      1
        70      1
        250      0
        ..
        124      1
        256      0
        265      0
        113      1
        185      0
        Name: target, Length: 61, dtype: int64
```

```
In [65]: 1 # Plot ROC curve and calculate AUC
        2 plot_roc_curve(gs_log_reg, x_test, y_test)
```


```
Out[65]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x136a84df340>
```



```
In [67]: 1 # Confusion Matrix
        2 print(confusion_matrix(y_test, y_preds))
```

```
[[17  8]
 [ 4 32]]
```

```
In [69]: ▶ 1 sns.set(font_scale=1.5)
2
3 def plot_conf_mat(y_test, y_preds):
4     """
5     Plots a nice looking confusion matrix using Seaborn's heatmap()
6     """
7
8     fig, ax = plt.subplots(figsize=(3,3))
9     ax = sns.heatmap(confusion_matrix(y_test, y_preds),
10                      annot=True,
11                      cbar=False)
12     plt.xlabel("True label")
13     plt.ylabel("Predicted label")
14
15 plot_conf_mat(y_test, y_preds)
```



now we've got a ROC curve, AUC metric and a confusion matrix, let's get a classification report, cross validated precision, recall and f1 score.

```
In [72]: ▶ 1 print(classification_report(y_test, y_preds))
```

	precision	recall	f1-score	support
0	0.81	0.68	0.74	25
1	0.80	0.89	0.84	36
accuracy			0.80	61
macro avg	0.80	0.78	0.79	61
weighted avg	0.80	0.80	0.80	61

Calculate evaluation metrics using cross validation

we're going to calculate accuracy, precision, recall and f1 score of our model using cross-validation and to do so we'll be using `cross_val_score()`

```
In [73]: ▶ 1 # Check best hyperparameters
2 gs_log_reg.best_params_
```

```
Out[73]: {'C': 0.1082636733874054, 'solver': 'liblinear'}
```

```
In [74]: 1 # Create a new classifier with best parameters
2 clf = LogisticRegression(C=0.1082636733874054, solver="liblinear")
3
4
```

```
In [76]: 1 # accuracy
2 cv_acc = cross_val_score(clf, x,y, cv=5, scoring="accuracy")
3 cv_acc.mean()
```

Out[76]: 0.834808743169399

```
In [78]: 1 # Preciaon
2 cv_precision = cross_val_score(clf, x,y, cv=5, scoring="precision")
3 cv_precision.mean()
```

Out[78]: 0.8182683982683983

```
In [79]: 1 # Recall
2 cv_recall = cross_val_score(clf, x,y, cv=5, scoring="recall")
3 cv_recall.mean()
4
```

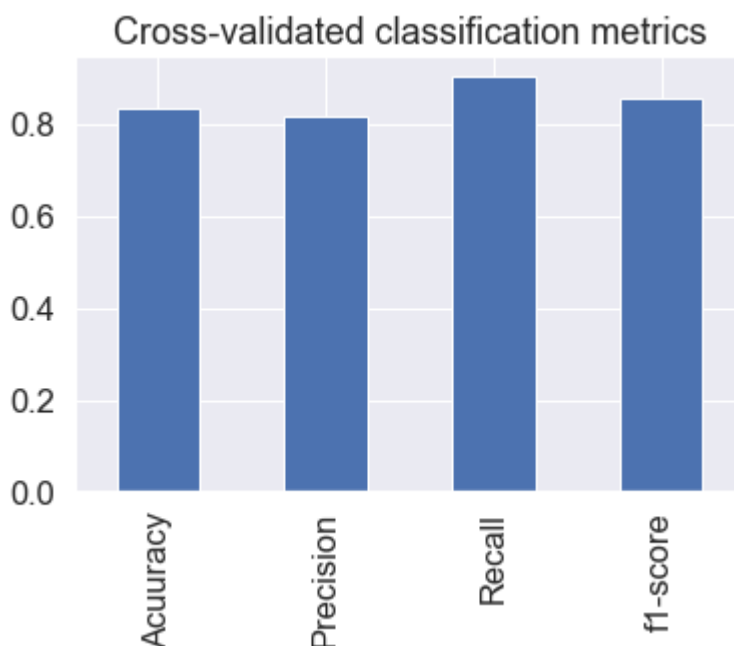
Out[79]: 0.9030303030303031

```
In [80]: 1 cv_f1 = cross_val_score(clf, x,y, cv=5, scoring="f1")
2 cv_f1.mean()
```

Out[80]: 0.8572876223964057

```
In [91]: 1 # Visualize cross-validated metrics
2 cv_metrics = pd.DataFrame({"Acuuracy": cv_acc.mean(),
3                             "Precision": cv_precision.mean(),
4                             "Recall": cv_recall.mean(),
5                             "f1-score":cv_f1.mean()}),
6                             index=[0])
7 cv_metrics.T.plot.bar(title = "Cross-validated classification metrics", legend=
```

Out[91]: <matplotlib.axes._subplots.AxesSubplot at 0x136a84524c0>



Feature Importance

Feature importance is another as asking, "Which feature contributed most to the outcomes of the model and how did they contributed?"

Finding feature importance is different for the each machine learning model. One way to find feature importance is ti search for (MODEL NAME) feture importance in google.

Let's find feature importance for our linear regression model.

```
In [93]: 1 # Fit an instance of LogisticRegression
2 clf = LogisticRegression(C=0.1082636733874054, solver="liblinear")
3
4 clf.fit(x_train, y_train);
```

```
In [96]: 1 df.head()
```

```
Out[96]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

```
In [94]: 1 # Check coef_
2 clf.coef_ #found after research
```


```
Out[94]: array([[ 0.01384033, -0.64080406,  0.60775315, -0.01685947, -0.00562062,
-0.07481902,  0.19251575,  0.03056174, -0.39713818, -0.44676097,
 0.21185815, -0.52018241, -0.5375065 ]])
```

```
In [97]: 1 # Match coef_ to features to coloms
2 feature_dict = dict(zip(df.columns, list(clf.coef_[0])))
3 feature_dict
```

```
Out[97]: {'age': 0.013840332662822355,
'sex': -0.6408040648369295,
'cp': 0.6077531458280921,
'trestbps': -0.016859471566922298,
'chol': -0.005620621169908997,
'fbs': -0.07481902221460926,
'restecg': 0.19251575074666608,
'thalach': 0.030561741114448537,
'exang': -0.3971381760345936,
'oldpeak': -0.44676097086868655,
'slope': 0.21185814998312974,
'ca': -0.5201824122684032,
'thal': -0.5375064986355229}
```

In [98]:

```
1 # Visualize feature importance
2 feature_df = pd.DataFrame(feature_dict, index=[0])
3 feature_df.T.plot.bar(title="Feature importance", legend=False);
4
```



In [99]:

```
1 pd.crosstab(df["sex"], df["target"])
```

Out[99]:

target	sex	
	0	1
0	24	72
1	114	93

In [100]:

```
1 pd.crosstab(df["slope"], df["target"])
```

Out[100]:

target	slope	
	0	1
0	12	9
1	91	49
2	35	107

6.Experimentation

If you haven't hit your evaluation metric yrt... ask yourself...

- Could you collect more data?

- Could you try a better model? Like CatBoost or XGBoost?
- Could you improve the current models? (beyond what we've done so far)
- If your model is good enough (you have hit your evaluation metric) we can save export model and share with others.

In []: ▶

1